

Report

In order to do the experiment that determines the best fit configuration, I have following methodology for optimization:

Use *cacheesim* to model and simulate different cache configurations with realistic workloads. Simulations can provide detailed insights into access times, miss rates, and power consumption.

Analyze access patterns of typical workloads(The given trace files) to understand the impact of B and S on hit and miss rates.

Calculate the overhead for different configurations. Highly associative caches require more tag bits per block, increasing the metadata storage.

Systematically vary B and S to find configurations that meet the optimization criteria. Pay special attention to points where further increases in B or S yield minimal improvements in access time or significantly increase metadata storage, area, or power.

The shell script I used to execute the program with different configurations is as follows:

```
1  #!/bin/bash
2
3  # Fixed parameters
4  c=15
5  C=17
6  trace_file="./traces/gcc.trace"
7  output="test.log"
8
9  # Iterate over the parameters within the specified ranges
10 for b in {5..7}; do
11     for s in {0..5}; do
12         for S in {1..6}; do # Start S greater than s to meet the restriction
13             for P in {0..2}; do
14                 for I in "LIP" "MIP"; do
15                     for r in "LFU" "LRU"; do
16                         # Ensure the size of the L2 cache is strictly
greater than the size of the L1 cache
17                         if ((C > c)) && ((S > s)) && ((C - S > c - s)); then
18                             ./run.sh -c "$c" -b "$b" -s "$s" -C "$C" -S "$S"
-P "$P" -I "$I" -r "$r" -f "$trace_file" >>"$output"
19                             fi
20                         done
21                     done
22                 done
23             done
24         done
25     done
26
```

The following python script is to find the smallest Average Access Time(AAT) for L1:

```
1  def find_smallest_l1_aat(file_path):
```

```

2     smallest_aat = None
3
4     with open(file_path, 'r') as file:
5         for line in file:
6             if line.startswith("L1 average access time (AAT):"):
7                 # Extract the numeric value of AAT
8                 aat = float(line.split(":")[1].strip())
9
10                # Initialize smallest_aat or compare with the current
smallest
11                if smallest_aat is None or aat < smallest_aat:
12                    smallest_aat = aat
13
14            return smallest_aat
15
16 # Replace 'test.log' with the actual path to your log file
17 file_path = 'test.log'
18 smallest_aat = find_smallest_l1_aat(file_path)
19 print(f"The smallest L1 average access time (AAT) is: {smallest_aat}")

```

File *gcc.trace* : The smallest average access time (AAT) for L1 among the provided configurations is 2.044 ns, which is achieved with the cache settings of:

```

1 Cache Settings
2 -----
3 L1 (C,B,S): (15,7,5). Replacement policy: LRU. Prefetcher disabled.
4 L2 (C,B,S): (17,7,6). Replacement policy: LRU. +1 prefetcher. Prefetch
insertion policy: LIP.
5
6 Cache Statistics
7 -----
8 Reads: 315500
9 Writes: 184500
10
11 L1 accesses: 500000
12 L1 hits: 495015
13 L1 misses: 4985
14 L1 hit ratio: 0.990
15 L1 miss ratio: 0.010
16 L1 average access time (AAT): 2.044
17
18 L2 reads: 4985
19 L2 writes: 2055
20 L2 read hits: 3818
21 L2 read misses: 1167
22 L2 prefetches: 840
23 L2 read hit ratio: 0.766
24 L2 read miss ratio: 0.234
25 L2 average access time (AAT): 29.510

```

- **L1 Cache** (C=15, B=7, S=5) : 14,336 bits
- **L2 Cache** (C=17, B=7, S=6) : 55,296 bits

File *leela.trace* : The smallest average access time (AAT) for L1 among the provided configurations is 1.799 ns, which is achieved with the cache settings of:

```

1  Cache Settings
2  -----
3  L1 (C,B,S): (15,7,4). Replacement policy: LRU. Prefetcher disabled.
4  L2 (C,B,S): (17,7,5). Replacement policy: LRU. +1 prefetcher. Prefetch
   insertion policy: LIP.
5
6  Cache Statistics
7  -----
8  Reads: 352008
9  Writes: 147992
10
11 L1 accesses: 500000
12 L1 hits: 499464
13 L1 misses: 536
14 L1 hit ratio: 0.999
15 L1 miss ratio: 0.001
16 L1 average access time (AAT): 1.799
17
18 L2 reads: 536
19 L2 writes: 14
20 L2 read hits: 323
21 L2 read misses: 213
22 L2 prefetches: 123
23 L2 read hit ratio: 0.603
24 L2 read miss ratio: 0.397
25 L2 average access time (AAT): 45.839

```

- **L1 Cache** (C=15, B=7, S=4): 14,080 bits
- **L2 Cache** (C=17, B=7, S=5): 54,272 bits

File *linpack.trace* : The smallest average access time (AAT) for L1 among the provided configurations is 4.136 ns, which is achieved with the cache settings of:

```

1  Cache Settings
2  -----
3  L1 (C,B,S): (15,7,3). Replacement policy: LRU. Prefetcher disabled.
4  L2 (C,B,S): (17,7,4). Replacement policy: LRU. +1 prefetcher. Prefetch
   insertion policy: LIP.
5
6  Cache Statistics
7  -----
8  Reads: 332993
9  Writes: 167007
10
11 L1 accesses: 500000
12 L1 hits: 478869
13 L1 misses: 21131
14 L1 hit ratio: 0.958
15 L1 miss ratio: 0.042
16 L1 average access time (AAT): 4.136
17
18 L2 reads: 21131
19 L2 writes: 20875
20 L2 read hits: 10490

```

```

21 | L2 read misses: 10641
22 | L2 prefetches: 10494
23 | L2 read hit ratio: 0.496
24 | L2 read miss ratio: 0.504
25 | L2 average access time (AAT): 56.457

```

- **L1 Cache** (C=15, B=7, S=3): 13,824 bits
- **L2 Cache** (C=17, B=7, S=4): 53,248 bits

File *matmul_naive.trace* : The smallest average access time (AAT) for L1 among the provided configurations is 3.321 ns, which is achieved with the cache settings of:

```

1 | Cache Settings
2 | -----
3 | L1 (C,B,S): (15,7,5). Replacement policy: LFU. Prefetcher disabled.
4 | L2 (C,B,S): (17,7,6). Replacement policy: LFU. +1 prefetcher. Prefetch
  | insertion policy: LIP.
5 |
6 | Cache Statistics
7 | -----
8 | Reads: 494452
9 | Writes: 5548
10 |
11 | L1 accesses: 500000
12 | L1 hits: 385860
13 | L1 misses: 114140
14 | L1 hit ratio: 0.772
15 | L1 miss ratio: 0.228
16 | L1 average access time (AAT): 3.321
17 |
18 | L2 reads: 114140
19 | L2 writes: 981
20 | L2 read hits: 113249
21 | L2 read misses: 891
22 | L2 prefetches: 808
23 | L2 read hit ratio: 0.992
24 | L2 read miss ratio: 0.008
25 | L2 average access time (AAT): 6.881

```

- **L1 Cache** (C=15, B=7, S=5): 14,336 bits
- **L2 Cache** (C=17, B=7, S=6): 55,296 bits

File *matmul_tiled.trace* : The smallest average access time (AAT) for L1 among the provided configurations is 1.909 ns, which is achieved with the cache settings of:

```

1 | Cache Settings
2 | -----
3 | L1 (C,B,S): (15,7,3). Replacement policy: LRU. Prefetcher disabled.
4 | L2 (C,B,S): (17,7,4). Replacement policy: LRU. +1 prefetcher. Prefetch
  | insertion policy: MIP.
5 |
6 | Cache Statistics
7 | -----
8 | Reads: 482548

```

```

9  writes: 17452
10
11  L1 accesses: 500000
12  L1 hits: 498182
13  L1 misses: 1818
14  L1 hit ratio: 0.996
15  L1 miss ratio: 0.004
16  L1 average access time (AAT): 1.909
17
18  L2 reads: 1818
19  L2 writes: 311
20  L2 read hits: 1135
21  L2 read misses: 683
22  L2 prefetches: 610
23  L2 read hit ratio: 0.624
24  L2 read miss ratio: 0.376
25  L2 average access time (AAT): 43.669

```

- **L1 Cache** (C=15, B=7, S=3): 13,824 bits
- **L2 Cache** (C=17, B=7, S=4): 53,248 bits

File *mcfl.trace* : The smallest average access time (AAT) for L1 among the provided configurations is 1.918 *ns*, which is achieved with the cache settings of:

```

1  Cache Settings
2  -----
3  L1 (C,B,S): (15,7,4). Replacement policy: LRU. Prefetcher disabled.
4  L2 (C,B,S): (17,7,5). Replacement policy: LRU. +1 prefetcher. Prefetch
   insertion policy: MIP.
5
6  Cache Statistics
7  -----
8  Reads: 304587
9  Writes: 195413
10
11  L1 accesses: 500000
12  L1 hits: 498969
13  L1 misses: 1031
14  L1 hit ratio: 0.998
15  L1 miss ratio: 0.002
16  L1 average access time (AAT): 1.918
17
18  L2 reads: 1031
19  L2 writes: 709
20  L2 read hits: 256
21  L2 read misses: 775
22  L2 prefetches: 766
23  L2 read hit ratio: 0.248
24  L2 read miss ratio: 0.752
25  L2 average access time (AAT): 81.270

```

- **L1 Cache** (C=15, B=7, S=4): 14,080 bits
- **L2 Cache** (C=17, B=7, S=5): 54,272 bits